

After the fan-connection era.

Question: Taylor Swift rose by connecting with fans; that playbook is saturated. What is the next paradigm, and where should LiveQue aim? Five research agents worked doctoral lenses (scene sociology, cultural economics, AI-scarcity economics, the parasocial-to-paid-participation shift, and forecasting with real prediction-market signals). A contrarian superforecaster then attacked the consensus, and the synthesis was rebuilt on what survived. Notably, the contrarian flagged and discounted client-flattery convergence before the final call.

THE PARADIGM

Retention replaces discovery: feeds — increasingly AI-curated — will still rent musicians their audience, but careers will be built by converting live co-presence into owned, off-algorithm patron relationships; in the post-Swift era you do not rise by being found, you rise by being kept.

Probability-weighted scenarios, 2027-2030

%	Scenario	What it looks like	Leading indicator
40	Algorithmic-curator succession (AI-mediated discovery)	The reference class is decisive: every discovery paradigm since 1981 (MTV, Napster, MySpace, YouTube, SoundCloud, playlists, TikTok) was a media-replication surface, never live infrastructure, and paradigm cadence has compressed to ~5 years with TikTok's window measurably closed. The successor is a feed — most plausibly AI-assistant/AI-DJ curation where breaking requires machine legibility (clean metadata, provenance, licensing-clear stems) rather than fan connection. Rising musicians optimize for the curator model, not the human scroll. Consistent with per-stream dilution: discovery value accrues to the algorithm owner, not the discovered.	Share of DSP streams initiated by AI assistants/DJ features crossing ~10%; any new app where unknown-artist snippets convert to saves at abnormal rates among 13-24s; first label signings publicly attributed to AI-curator traction.
27	No new paradigm: fragmentation grind	Fragmentation IS the observed regime since ~2022 and deserves far more than the consensus 10%: A&R screens ~20 channels, first-time-artist Hot 100 top-10 debuts sit at millennium lows, and no single surface has replaced TikTok in four years. Breaking remains multi-channel attrition plus touring; the winners are artists who treat every channel as rented and build direct income early. Functionally this scenario converges with the paradigm sentence — when no discovery lottery exists, retention economics are all there is.	First-time-artist top-10 debut rate staying at/near lows through 2027; A&R channel-fragmentation persisting in trade reporting; no app achieving TikTok-2019-style unknown-artist breakout rates for eight consecutive quarters.

%	Scenario	What it looks like	Leading indicator
20	Superfan patronage deepening (top-heavy, tooling trickles down)	The industry's own growth model: Goldman's \$4.3B superfan uplift, expanded-rights/D2C growing fastest (+21.5%, MIDiA), Luminate superfans spending \$113/mo on live. The contrarian caveat holds — the mechanism is engineered at the idol tier (Weverse, K-pop) and Gini-0.829 math means spend concentrates upward, so this is not automatically a working-musician paradigm. But the TOOLING (super-premium tiers, D2C rails, patron CRM) commoditizes downward, and patrons demonstrably overpay in reciprocity relationships (Regner 2021), which is exactly the room-scale transaction. Bandcamp Fridays (\$154M cumulative, record 2025) prove '0% take' is a durable demand-side trigger.	DSP super-premium tiers shipping and holding subscribers; expanded-rights/D2C segment growth staying >15% YoY in MIDiA data; patron-CRM tools appearing at sub-\$20/mo price points for independent artists.
13	Verified-live / proof-of-human credential	Real but overweighted at the original 35%. The supply flood is uncontested (AI >50% of Deezer daily uploads; 97% listener indistinguishability) and provenance infrastructure is genuinely shipping (Deezer tagging live since June 2026; Spotify/DDEX disclosure fields). But the credential currently has no buyer — A&R practice is explicitly digital-first, badges verify recordings not gigs, and tip-timestamp attestations are as gameable as streams (85% of AI streams are already fraud). This becomes the paradigm only if a certification regime extends to performance AND the signing chain starts consuming live-demand data. Cheap to hedge, wrong to bet on.	A major DSP or DDEX shipping a human-verified badge that covers PERFORMANCE (not just recording provenance); first publicized label signing attributed to live-sourced demand data. Estimated ~30-35% chance one occurs by end-2028 — not the consensus 55%.

Where LiveQueue actually stands (the honest part)

Honest answer: LiveQueue is not positioned to own the most probable winning scenario, and cannot be. The 40% scenario (AI-curated discovery) and the 27% scenario (fragmentation) are both feed phenomena; no live-infrastructure company has ever owned a discovery paradigm because rooms do not scale, and the five-lens convergence on 'your product seeds the next paradigm' was narrative gravity — five lenses drawing on the same 2024-26 think-piece corpus is one correlated signal, not consilience. But the question conflates two things the data cleanly separates: how musicians RISE (still media replication — LiveQueue has no purchase there) and how musicians EARN AND SURVIVE (genuinely shifting live and direct — LiveQueue sits exactly there). The earning side strengthens under ALL four scenarios: working musicians already take ~81% of music income live (MIRA), per-stream dilution is arithmetic, cash tips are structurally dying (14% of US payments and falling), and the fan relationship captured at the moment of peak arousal — the tipped request — is the one asset that compounds whether the next feed is AI-curated, fragmented, or never arrives. Sober caveats that bound the ambition: grassroots demand is price-weak (MVT: attendance +13%, value down, 53% of venues at zero profit); digital-tip fatigue is measurable (63% of Americans negative on tipping, guilt-tip spend -38%) even if performer tipping carries more goodwill than counter-service; and 0% rake is an adoption wedge, not a moat — Sofar survived its scandal and Bandcamp's fair-trade economics didn't save it, so revenue must eventually come honestly from performer-side SaaS. The realistic prize: become the default request-tip-CRM utility for the

working-musician tier — an underserved niche where every surviving fact helps and no paradigm needs to arrive — while holding a near-free option on the 13% verified-live scenario. For a two-person, zero-budget team, the base rate for 'niche utility becomes paradigm owner' is ~zero; the base rate for 'well-loved niche utility sustains its founders' is decent. That asymmetry is the whole decision.

The three strategic focuses, next 18 months

1. Make every tip mint a performer-owned fan record (the tip is the transaction; the relationship is the product)

Why: This is the single strategy that pays in every scenario, including the 27% grind: the conversion event from listener to patron happens at high-arousal live moments (peer-reviewed tipping and patronage literature), Stripe already captures contact at that moment, and an owned, off-algorithm fan list is the one compounding asset the working-musician tier lacks. It upgrades LiveQue from gratuity widget to the CRM of the room — and it is what makes the open-mic density plan accrete into something instead of resetting every Tuesday: each saturated room deposits permanent fan relationships, not just tips.

First step: Add an opt-in 'follow this artist / tell me their next gig' checkbox to the tip-confirmation screen, persist it to a per-performer fan list with export, and ship one re-engagement loop (next-gig notification). Weeks of work on existing Stripe rails; validate open-in rates with the current LA cohort before building anything further.

2. Solve revenue honestly: 0% on tips forever, paid performer tier for the tools

Why: The contrarian finding that survives hardest: zero-rake is a genuine demand-side trigger (Bandcamp Fridays' \$154M proves fans treat 'artist keeps 100%' as a purchase reason) but forgoing all revenue is the absence of a business model, not a moat — any funded entrant can match 0% on tips and outspend you. The defensible position is 0% on the money plus charging performers modestly for the compounding layer (fan CRM, setlist/request analytics, receipts page, branding). This preserves the wedge, funds the team, and cannot be attacked reputationally because the fan-to-artist dollar is never touched.

First step: Instrument feature usage now; when the fan-list feature from focus 1 shows retention, price a \$5-10/mo Pro tier and hand-sell it to the 5-10 most active open-mic performers before writing any billing code. Target: 20 paying performers by month 12 as proof the tier clears.

3. Persist the exhaust as a cheap option on verified-live — a hedge, not a bet

Why: The 13% scenario is the only one where LiveQue could own something structural (the grassroots system of record for human liveness), and the option costs almost nothing: it is a schema decision (never delete request/tip/song/venue/timestamp data) plus one shareable page. Do NOT architect the company around certification, demand graphs, or DSP APIs — there is currently no buyer for the credential and the attestation is gameable. But if the tripwire fires (a performance-covering human-verified badge, or a publicized signing off live-demand data), being the incumbent holder of years of per-room, per-dollar demand records is a position no competitor can backfill, because it only accrues night by night.

First step: Make completed requests permanent records rather than transient queue entries, and ship a public per-performer 'live receipts' page (shows played, songs performed, count of real paying requesters — amounts anonymized) designed to be linked from a Spotify/Instagram bio. One schema migration plus one page; then stop investing until the tripwire fires.

What the contrarian killed or weakened

First, a methodological red flag: five nominally independent doctoral lenses all converged on 'the client's product is the seed of the next paradigm.' That is not consilience — it is narrative gravity plus client-flattery. All five draw from the same 2024-26 think-piece ecosystem (live-is-back, AI-flood, superfan extraction), so their agreement is one correlated signal, not five independent ones. Discount the convergence itself. WHAT SURVIVES ADVERSARIAL REVIEW: (1) AI is commoditizing recorded audio — supply fact, uncontested. (2) Working musicians earn overwhelmingly live and direct, and per-stream dilution is arithmetic — robust. (3) Cash-to-digital tip substitution is mechanical: the marginal tip must go digital. (4) Grassroots co-presence demand exists but is PRICE-WEAK — MVT 2025 shows attendance up 13% while ticket value falls and 53% of venues make zero profit. Together these support a modest thesis, not a paradigm thesis. WHAT FAILS: the structural IRL-revival tailwind (ATUS through 2025 shows socializing still declining, worst among the young — the revival is a niche counter-current inside a secular decline); the proof-of-liveness credential (no buyer — A&R remains digital-first, badges verify recordings not gigs, and tip timestamps are as gameable as streams); prediction-market backing (thin, and music markets mainly manufacture fraud); room-scale Superchat economics (decades-old TouchTunes/piano-bar niche that never minted a career; winner-take-most math plus tip fatigue); and zero-rake-as-moat (Sofar and Bandcamp jointly disprove both the punishment mechanism and the durability of fair-trade economics). THE CENTRAL CONFLATION: the lenses blur how musicians EARN (genuinely shifting live and direct — survives) with how musicians RISE (still media replication — no evidence of a shift, because rooms don't scale and every paradigm since MTV has been a feed). LiveQue sits on the earning side. That is fine — sell to earners. RECALIBRATED SCENARIO VECTOR, 2027-2030: feed/algorithmic successor including AI-curated discovery ~40-45%; no-new-paradigm slow grind ~25-30% (fragmentation IS the current regime, so this deserves far more than the consensus's 10%); superfan financialization, concentrated at the idol tier ~15%; verified-live/proof-of-human as THE breakout paradigm ~8-12%. Tripwire probability (major DSP human-verified badge OR publicized label signing off live-demand data by end-2028): ~30-35%, not 55%. WHAT LIVEQUE SHOULD ACTUALLY DO: the calibrated prize is unglamorous and real — become the default gratuity-and-request utility for working musicians, a genuinely underserved niche where every surviving fact helps and no paradigm needs to arrive. Keep the LA open-mic saturation plan (it is sound bounded-community diffusion regardless of the Centola dressing). The single best idea in the consensus is also its least glamorous: capture the fan contact at the moment of the tip so each

performer accretes an owned fan list — that pays off in EVERY scenario including slow-grind, and is just CRM logic. Persist the request/tip exhaust because it costs nearly nothing to keep — a cheap hedge on the verified-live scenario, not a bet on it. Keep 0% on tips as the adoption wedge but solve revenue honestly (performer-side SaaS), because forgoing all revenue is an absence of a business model, not a moat. Do NOT architect the company around certification layers, demand graphs, or paradigm ownership: the base rate for 'niche utility becomes paradigm owner' is approximately zero, while the base rate for 'well-loved niche utility sustains a two-person zero-budget team' is decent — and for these founders, that asymmetry is the whole decision. Sources consulted: [Axios/ATUS 2026](<https://www.axios.com/2026/07/05/americans-socializing-decline>), [MVT Annual Report 2025](<https://www.musicvenue trust.com/2026/01/music-venue-trust-annual-report-2025/>), [Bankrate tipping survey](<https://www.bankrate.com/credit-cards/news/tipping-culture-survey/>), [Popmenu via PRNewswire](<https://www.prnewswire.com/news-releases/77-of-consumers-say-tipping-in-the-us-has-become-ridiculous-according-to-popmenus-annual-study-302577741.html>), [CNBC on thin prediction markets](<https://www.cnbc.com/2026/07/02/prediction-markets-mostly-have-thinly-traded-contracts-.html>), [Rolling Stone on Velvet Sundown](<https://www.rollingstone.com/music/music-features/velvet-sundown-ai-band-suno-1235377652/>), [Influor A&R playbook](<https://www.influor.com/blog/the-new-a-r-playbook-what-record-labels-need-to-know-about-signing-in-the-creator-economy>).

COLLAPSES — Lens 1 (scenes/third places): 'Demand for co-presence is a structural public-health correction, not a trend — a demand tailwind for live gathering through 2030.'

The best time-use data says the opposite is still happening. American Time Use Survey data through 2025 shows face-to-face socializing falling from 47 to 35 min/day since 2003, with 15-24-year-olds — the cohort the IRL-revival story is about — down 43% (61 to 35 min/day), and time attending/hosting social events down ~50-69% since 2003 (Axios, July 2026; 'The Anti-Social Century' corpus). The loneliness epidemic has produced awareness and think-pieces, not a behavioral reversal: revealed preference keeps shifting toward screens every single year the claim says a correction is underway. Run clubs and coffee raves are real, growing niches — but they are small absolute counter-currents inside a large secular decline, and the 'IRL revival' narrative is produced by the same media ecosystem that generated the decline coverage. A structural tailwind that never appears in the aggregate behavioral data is a vibe, not a tailwind.

WEAKENED — Lens 1: 'Scissors crisis — rising gathering demand meets contracting venue stock, so music migrates to bar-less spaces where direct audience-to-artist payment replaces the bar cut.'

Music Venue Trust's 2025 annual report complicates this badly: UK grassroots attendance actually rose 13% (21M attendances) — but ticket VALUE fell, 53% of venues made zero profit, margins for the profitable minority were 2.5%, 6,000 jobs were cut, and 30 venues closed. So co-presence demand at this tier exists but is price-weak: people show up and pay less. That is the opposite of the monetization inference. The migration to free/cheap formats (parks, run clubs, coffee raves) signals declining willingness-to-pay at exactly the tier LiveQue serves, and free-entry formats historically monetize at busking economics — cents per head. The migration is real; the assumption that direct fan-to-artist payment will replace the bar cut at meaningful scale is asserted, not evidenced.

WEAKENED — Lens 1: 'Zero-rake is the intimate economy's license to operate and a moat — the Sofar/Groupmuse natural experiment shows violators are punished reputationally.'

Sofar Sounds is the direct counterexample to its own citation: repeatedly exposed since 2019 for paying artists ~\$100 flat on shows grossing far more, settling with New York regulators over labor practices — and it kept operating and expanding globally. Live Nation runs the most-resented fee structure in music and posted record FY2025 results. Reputational punishment in live music is empirically toothless. Meanwhile Bandcamp — the actual fair-trade exemplar — was sold twice and gutted in 2023-24: artist-friendly economics attract goodwill but do not defend themselves or the company. 0% is a genuine adoption wedge, but a moat that consists of forgoing all revenue is also the absence of a business model; any funded competitor can match 0% on tips and monetize elsewhere (SaaS, data, payments float).

WEAKENED — Lens 1 + Lens 4: 'Growth is concentrated in participatory formats (jams, open mics, run clubs), and a visible pay-to-steer ritual can monetize that porous audience/performer boundary.'

Turino — cited in support — argues the reverse: participatory music RESISTS commodification; its value derives from sitting outside market exchange. Motivation-crowding research (Frey & Oberholzer-Gee; Gneezy & Rustichini; the Titmuss tradition) predicts that inserting visible price-ranking into a communal ritual crowds out the intrinsic motives that make the ritual work. Participation inequality (the 90-9-1 pattern) means most of any room never transacts. And 2025-26 tipping data shows visible payment prompts now generate measurable resentment: 66% of consumers feel pressured by digital tip screens, and 'guilt tipping' spend fell 38% year-over-year (Popmenu 2025). A tip-ranked public queue projected over a communal jam is the highest-pressure version of the prompt Americans say they hate. The formats are growing precisely because they are cheap and non-transactional.

WEAKENED — Lens 2: 'Digital prompts add a measured 15-26% tipping premium; the marginal tip goes digital as cash floors out.'

The substitution half is mechanical and survives — cash share of payments genuinely floors around 14%, so tip capture must go digital. But the premium half rests on novelty-window field experiments (e.g., Edinburgh Fringe first-exposure effects) while population-level data shows the curve bending the other way under repeated exposure: 63% of Americans now hold negative views of tipping (up from 59% YoY, Bankrate 2025), average tip rates dipped below 15% (Square 2025), the share tipping 20%+ at restaurants fell 37% to 35%, and 77% call US tipping culture 'ridiculous' (Popmenu). Tip-prompt proliferation is producing habituation and backlash, not a durable premium. Musician tipping enjoys more goodwill than counter-service tipping, so some premium likely persists — but banking on 15-26% through 2030 extrapolates a honeymoon.

WEAKENED — Lens 2: 'Superfan monetization (Goldman \$4.3B; Luminate \$113/month live spend) is the industry growth engine LiveQue can ride at the working-musician tier.'

Composition fallacy. Goldman's figure is a sell-side model of superstar-fandom extraction (deluxe tiers, VIP packages, K-pop-style apps); Luminate's superfan spends \$113/month largely on premium tickets to name acts. Where superfan monetization demonstrably works — Weverse, K-pop — it runs on industrially produced parasocial idol intimacy, the exact mechanism Lens 4 declares saturated (an unresolved contradiction across the consensus). The lens's own MIRA data shows the median working musician earns \$21.3K, i.e., the tier is poor: superfans per median artist may be single digits, and the winner-take-most math cited elsewhere (Gini 0.829) predicts superfan spend concentrates at the top exactly as streams did. Industry-aggregate superfan growth does not distribute to the bar-gig tier; there is no evidence it ever has.

COLLAPSES — Lens 3: 'Verified-live-human performance becomes the scarce premium good, and a tip-backed performance ledger becomes the grassroots artist's portable certified-human credential.'

Three independent failures. (a) Authenticity premiums historically decay with habituation: autotune, drum machines, playback DJs, and ghost production each triggered panics that faded; The Velvet Sundown retained ~1M monthly listeners through and after its July 2025 AI reveal (Rolling Stone, CBC) — revealed indifference, not punishment — and lab attribution effects (music rated lower when labeled AI) are small, novelty-loaded, and demand-characteristic-prone. (b) No buyer for the credential: 2025-26 A&R practice is explicitly digital-first — scouts screen on TikTok/streaming momentum thresholds, and industry guidance is blunt that a creator with 800K engaged followers is a lower-risk signing than an artist with a decade of touring and no digital footprint. Nobody in the signing chain requests open-mic receipts, DSP provenance badges verify recordings not gigs, and at grassroots nobody doubts the performer is human — they can see them. Proof-of-liveness solves a non-problem at the only tier LiveQue occupies. (c) The attestation is weak: a tip timestamp is trivially manufactured by friends tipping friends — the same fraud pattern (85% of AI streams fraudulent) the lens says invalidates streaming data. The recorded-music supply flood is real; the certification-layer implication does not follow from it.

WEAKENED — Lens 4: 'The Twitch-native cohort will treat paying to steer a live room as normal; LiveQue should be the room-scale Superchat and own the interaction grammar.'

Paid song-steering in physical rooms already exists at scale and has for decades — piano-bar tip jars, wedding-band requests, TouchTunes at 12M paid plays/week — and in all that time it minted zero careers and zero paradigms; the lens's own existence proof shows it is a stable gratuity niche, not an ascendancy mechanism. The livestream evidence points the wrong way: Superchat's top earners are VTubers and idols — maximal parasocial constructs — contradicting both the 'ordinariness drives payment' premise and the lens's own 'parasocial is saturated' premise. The winner-take-most math it cites (top 1% capture 38%; Gini 0.829) means a 0%-rake layer across thousands of median performers aggregates trivial GMV and zero revenue — 'own the interaction layer' is only a prize if the layer can charge someone. And the cohort supposedly aging into bars is measurably going out less (15-24 socializing down 43% since 2003, still declining through 2025). Plausible as a UX insight; fails as an economic paradigm.

COLLAPSES — Lens 5: 'Prediction-market signals support the scenario vector — 35% verified-live paradigm, 55% tripwire by end-2028.'

The markets cited cannot bear this weight. CNBC's 2026 analysis found ~70% of closed Polymarket markets did under \$10K lifetime volume; music markets are a rounding error of the \$44.8B/month category total; and the flagship music-market story of 2026 was Spotify demanding takedowns from Kalshi/Polymarket after ~\$3M of chart-market volume induced 500K+ fraudulent streams — evidence that financializing music outcomes manufactures fraud, not signal. No market anywhere prices 'verified-live becomes the breakout paradigm'; the 35/22/18/15/10 vector is the analyst's priors wearing a market costume, with false precision. The reference-class argument also inverts the ordering: every discovery paradigm since 1981 — MTV, radio consolidation, Napster, YouTube, blogs, SoundCloud, TikTok — was a media-replication surface; zero were live-infrastructure paradigms, because rooms do not scale. The outside view puts most probability on another feed (human- or AI-curated) and far more than 10% on 'no new paradigm' — fragmentation has been the observed regime since ~2022.

SURVIVES — Lens 2 (stress-testing the strongest leg): 'Working musicians earn ~81% of music income live; live/direct share grows as recorded audio commoditizes.'

Attempted disconfirmation largely fails. The MIRA survey is 2018 and self-selected — but post-pandemic replications (UK Musicians' Union and subsequent artist-income surveys) reproduce live dominance for the non-star tier; streaming payout arithmetic (catalog growing faster than subscription revenue) guarantees continued per-stream dilution regardless of which paradigm arrives; and the AI supply flood into recorded audio is uncontested fact confirmed across Deezer's own disclosures. The one caveat from MVT 2025 is that grassroots live demand is price-weak (attendance up, spend down), so 'live income dominates' does not mean 'live income grows much.' This is the leg of the consensus that survives adversarial review.

Quarterly watch list

- Own funnel first (trumps all secondary data): per-room scan-to-request and request-to-tip conversion, repeat-tipper rate, and fan-list opt-in rate — if opt-in at the tip moment runs below ~15-20%, the retention thesis needs redesign, not more rooms.
- Tipping-fatigue series (Bankrate annual, Popmenu, Square tip-rate data): does the counter-service backlash (63% negative, guilt-tip spend -38% YoY) spread to performer tipping, or does the musician exemption hold? Falsifies the digital-tip premium assumption.
- Time-use data (ATUS annual release, plus Eventbrite-type intent surveys): does in-person socializing among 15-34s finally flatten or reverse? Through 2025 it is still declining — the 'IRL revival' remains a niche counter-current, so treat any venue-less/daytime-format expansion (coffee raves, run clubs) as opportunistic, not structural, until this line bends.
- AI-curation share (scenario 1 tripwire): streams initiated via AI DJ/assistant features in DSP disclosures; any new short-video or agent app showing abnormal unknown-artist save/breakout rates among 13-24s — if this fires, double down on retention tooling because discovery just got more algorithmic, not less.
- Provenance regime (scenario 4 tripwire): Deezer AI-tagging scope, Spotify/DDEX disclosure-field rollout — specifically whether any 'human-verified' construct extends to live PERFORMANCE rather than recordings; plus any publicized signing attributed to live-demand data. Either event escalates focus 3 from hedge to build.
- Earning-side confirmation (MIDiA quarterly/annual): expanded-rights/D2C segment growth staying >15% YoY and per-stream value continuing to dilute — the arithmetic backbone of the whole strategy; if D2C growth stalls two consecutive readings, revisit.
- Grassroots price-weakness (MVT annual report and any US analogue): attendance-up/spend-down divergence — if grassroots willingness-to-pay recovers, raise suggested tip defaults; if it worsens, the SaaS tier (focus 2) matters more than tip volume.

- Superfan tooling commoditization (scenario 3): DSP super-premium tier launches and retention, and the price floor of independent-artist patron-CRM tools — if credible competitors ship fan-CRM at <\$20/mo, LiveQue's differentiation must be the in-room capture moment, which no remote tool can replicate.

Key sources

MIRA/Princeton Survey Research Center, Survey of Musicians (2018) — 81% of working-musician income is live; median \$21.3K (psrc.princeton.edu)

MIDiA Research 2025/2026 — expanded rights/D2C fastest-growing segment +21.5%; per-stream dilution arithmetic (midiaresearch.com)

Goldman Sachs 'Music in the Air' 2025 — superfan uplift model, live to \$67.1B by 2035, recorded forecast cut (goldmansachs.com; via MBW)

Luminate superfan reports 2024-2026 — \$113/month live spend, superfan demographics (luminatedata.com)

Deezer Newsroom telemetry Jan 2025-Jul 2026 — AI >50% of daily uploads but 1-3% of streams, 85% of AI streams fraudulent; Deezer/Ipsos n=9,000: 97% cannot detect AI music (newsroom-deezer.com)

Federal Reserve, 2025 Diary of Consumer Payment Choice — cash at 14% of payments; digital-tip substitution is mechanical (frb.services.org)

'QR code-enabled tips to street performers at the Edinburgh Fringe,' Journal of Cultural Economics (2025) — signage/moment design is the conversion bottleneck (springer.com)

Music Venue Trust Annual Report 2025 — attendance +13% but 53% of venues at zero profit: grassroots demand is price-weak (musicvenuetrust.com)

American Time Use Survey via Axios (July 2026) — in-person socializing still declining, worst among 15-24s; falsifies the structural IRL-revival premise

Bankrate 2025 tipping survey and Popmenu annual study — 63% negative on tipping, guilt-tip spend -38% YoY: digital-tip premium is a honeymoon, not a constant

'Rich-Get-Richer? Platform Attention and Earnings Inequality using Patreon Earnings Data,' Industrial and Corporate Change (2026); Regner, Journal of Cultural Economics (2021) — patronage holds its median best off-algorithm; patrons voluntarily overpay

Bandcamp Fridays cumulative payout data (\$154M since 2020, record \$19M in 2025) — '0% take' as measured demand trigger (MBW/Billboard, Dec 2025)

CHI 2025 Superchat inequality study — Gini 0.829, top 1% take 38%: paid-participation economics are winner-take-most; plan for whales, monetize the layer not the lottery

Rosen, 'The Economics of Superstars,' AER (1981); Baumol & Bowen (1966) — the theoretical frame: streaming perfects winner-take-all while live scarcity structurally appreciates

TechCrunch on Sofar Sounds (2019) and Bandcamp's 2023-24 sale/gutting — joint disproof of reputational punishment and of fair-trade economics as a self-defending moat